

CMPS6610Lab03

Name (Team Member 1):_Orville Cunningham_____

Name (Team Member 2):_____

In this lab you will practice using the map and reduce functions. These functions are commonly used together in a map-reduce framework, used by Google and others to parallelize and scale common computations.

Part 1: Counting Words

In the first part, we will use map-reduce to count how often each word appears in a sequence of documents. E.g. if the input is two documents: ['i am sam i am', 'sam is ham']

then the output should be

```
[('am', 2), ('ham', 1), ('i', 2), ('is', 1), ('sam', 2)]
```

We have given you the implementation of the main map-reduce logic **def**

```
run_map_reduce(map_f, reduce_f, docs)
```

To use this function to count words, you'll need to implement your own `map_f` and `reduce_f` functions, described below.

1. Complete `word_count_map` and test it with `test_word_count_map`.
2. Complete `word_count_reduce` and test it with `test_word_count_reduce`.
3. If the above are correct, then you should now be able to test it the full solution `test_word_count`
4. Assume that a word `w` appears `n` times. What is the **work** and **span** of `word_count_reduce` for this word, assuming a parallel implementation of the reduce function?

Enter answer here.

The `word_count_reduce` function should take a word and a list of counts as input and return a tuple with the word and the sum of the counts. This function will be called for each word, and the list of counts will be all the 1s returned by the `word_count_map` function for that word. Here's an example implementation

5. Why are we going through all this trouble? Couldn't I just use this function to count words?

```
docs = ['i am sam i am', 'sam is ham']
```

```
counts = {} for doc
```

```
in docs:
```

```
    for term in doc.split():
```

```
        counts[term] = counts.get(term, 0) + 1
```

```
# counts = {'i': 2, 'am': 2, 'sam': 2, 'is': 1, 'ham': 1}
```

What is the problem that prevents us from easily parallelizing this solution?

5.The problem with the solution that uses a loop and a dictionary to count words is that it cannot be easily parallelized. This is because the updates to the dictionary cannot be safely performed by multiple threads or processes at the same time, as it would lead to race conditions and incorrect results.

Map-reduce is a way to solve this problem by splitting the data into chunks that can be processed in parallel, then combining the results.

Enter answer here

Part 2: Sentiment analysis

Finally, we'll adapt our approach above to perform a simple type of sentiment analysis. Given a document, rather than counting words, we will instead count the number of positive and negative terms in the document, given a predefined list of terms. E.g., if the input sentence is it was a terrible waste of time and the terms terrible and waste are in our list of negative terms, then the output is [('negative', 1), ('negative', 1)]

6. Complete the sentiment_map function to implement the above idea and test it with test_sentiment_map.

1

7. Since the output here is similar to the word count problem, we will reuse word_count_reduce to compute the total number of positive and negative terms in a sequence of documents. Confirm your results work by running test_sentiment.

